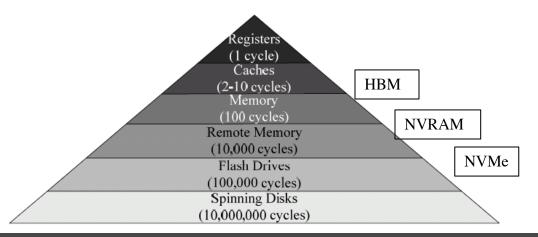


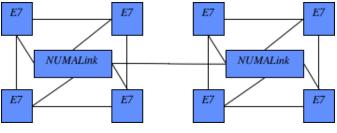


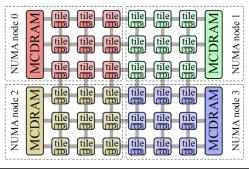
Current Supercomputer Architectures

- Multi-socket server nodes
 - NUMA
 - Accelerators
- High performance interconnect
 - e.g. Infiniband, OmniPath
- Scalable parallel approach needed to achieve performance



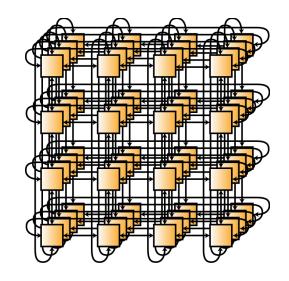


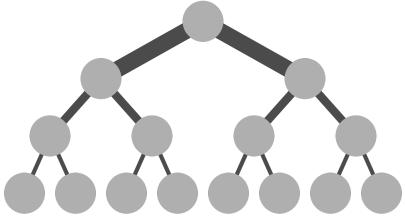




Network Topologies

- Mesh, Torus, Hypercube
- Tree based
 - Fat-tree
 - Clos
- Dragonfly
- Metrics
 - Bandwidth
 - Diameter, Connectivity
 - Bisection bandwidth





Parallel Computing

- Executing instructions concurrently on physical resources (not time slicing)
 - Multiple tightly coupled resources (e.g. cores) collaboratively solving a single problem

Benefits

- Capacity
 - · Memory, storage
- Performance
 - More instructions per unit of time (FLOPS)
 - Data streaming capability

Cost and Complexity

- Coordinate tasks and resources
- Use resources efficiently



Memory, Communication, and Execution Models

Shared

Communication model: shared memory

Distributed

Communication model: exchange messages

Execution Models

Fork-Join (e.g. Thread Level Parallelism)

Parallelism enabled by decomposing work

- Tasks can be executed concurrently
- Some tasks can have dependencies



What is OpenMP?

High level parallelism abstraction based on thread

- Easy to use
- Suitable for an incremental approach

A specification and evolving standard

- "a portable, scalable model ... for developing portable parallel programs"
- http://openmp.org
- GNU, Intel, PGI, etc.

A set of

- Compiler directives
- Library routines
- Environment variables
- Supports C/C++ and Fortran

```
#pragma omp parallel {
....
}
```

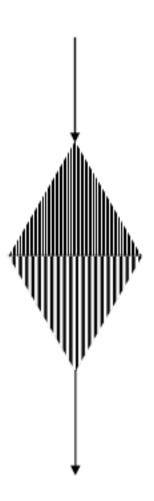
OpenMP Models

Fork/Join Execution

- Process starts single threaded (master thread)
- Forks child threads activated in parallel regions (team)
- The team synchronizes and threads are disbanded
 - barrier
- Overhead is mitigated by reusing threads
- Master thread continues execution of serial phases

Work decomposition

- Programming constructs
- Scope and compound statements
- Declarative in loops
- Mapping to threads can be static or dynamic
- Barriers and synchronization automatically inserted



Compiler Directives

- Compiler directives are the main mechanisms for introducing parallelism. Functionality enabled includes:
 - Spawning a parallel region
 - Diving code among threads
 - Distributing loop iterations over threads
 - Serialization of parts of the code
 - Synchronization of work
- Example:

#pragma omp parallel default(shared) private(beta,pi)



Parallel Region Construct

```
!$OMP PARALLEL [clause ...]

IF (scalar_logical_expression)

PRIVATE (list)

SHARED (list)

DEFAULT (PRIVATE | FIRSTPRIVATE | SHARED | NONE)

FIRSTPRIVATE (list)

REDUCTION (operator: list)

COPYIN (list)

NUM_THREADS (scalar-integer-expression)
```

code block
!\$OMP END PARALLEL



Data Scope Attribute Clauses

- PRIVATE variables in the list are private to each thread.
- SHARED variables in the list are shared between all threads.
- DEFAULT default scope for all variables in a parallel region.
- FIRSTPRIVATE variables are private and initialized according to value prior to entry into parallel or work sharing construct.
- LASTPRIVATE variables are private, the value from the last iteration or section is copied to original variable object.
- Others COPYIN, COPYPRIVATE
- REDUCTION reduction on variables in the list

Simple C Example

```
#include <omp.h>
int main(void) {
   int number_of_threads, thread_id;
   #pragma omp parallel private(thread_id)
      thread_id = omp_get_thread_num();
      #pragma omp master
         number_of_threads = omp_get_num_threads();
      #pragma omp barrier
      printf("hello, shared memory world from thread %d of \
         %d\n", thread_id, number_of_threads);
   return 0;
```

Simple FORTRAN Example

```
program hello_omp
     use, intrinsic :: iso_fortran_env
      implicit none
      integer :: number_of_threads, thread_id
      integer :: omp_get_num_threads, omp_get_thread_num
!$omp parallel private(thread_id)
     thread_id = omp_get_thread_num()
!$omp master
     number_of_threads = omp_get_num_threads()
!$omp end master
!$omp barrier
     write(unit=output_unit, fmt=*) "hello, shared memory&
        & world from thread ", thread_id, " of ",&
        & number of threads
!$omp end parallel
      stop
      end program hello_omp
```



Compiling on Expanse

- Load a compiler module. For example: module load gcc/10.2.0
- Compile commands:
 - Fortran: gfortran -fopenmp -o hello_fortran hello-omp.f90
 - C: gcc -fopenmp -o hello_c hello-omp.c

Sample job script

```
#!/bin/bash
#SBATCH --job-name="hello"
#SBATCH --output="hello.%j.%N.out"
#SBATCH --partition=shared
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=8
#SBATCH -account=XYZ123
#SBATCH --export=ALL
#SBATCH -t 00:10:00
module load gcc
#Export the number of OpenMP threads
export OMP NUM THREADS=8
#Run the openmp code
./hello fortran
./hello_c
```



Work-Sharing

Schedule:

- Static Loop iterations are statically divided (chunk or as close to even as possible)
- Dynamic Loop iterations are divided in size chunk, and dynamically scheduled among threads. When a thread finishes one chunk it is dynamically assigned another
- Guided Similar to dynamic but chunk size is proportionally reduced based on work remaining.
- Runtime set at runtime by environment variables
- Auto set by compiler or runtime system.



Number of Threads

- Number of threads will be determined in the following order of precedence:
 - Evaluation of the IF clause
 - Setting of NUM_THREADS clause
 - omp_set_num_threads() library function
 - OMP_NUM_THREADS environment variable
 - Default usually ends up being the *number of cores on the node* (!)
- The last factor can accidentally lead to oversubscription of nodes in hybrid MPI/OpenMP codes.



Simple OpenMP Program – Compute Pl

- Find the number of tasks and taskids (omp_get_num_threads, omp_get_thread_num)
- PI is calculated using an integral. The number of intervals used for the integration is fixed at 128000.
- Use OpenMP loop parallelization to divide up the compute work.
- Introduce concept of private and shared variables.
- OpenMP reduction operation used to compute the sum for the final integral.



OpenMP Program to Compute Pl

```
#include <omp.h>
#include <stdio.h>
#include <stdlib.h>
int main (int argc, char *argv[])
int nthreads, tid;
int i, INTERVALS;
double n 1, x, pi = 0.0;
INTERVALS=128000:
/* Fork a team of threads giving them their own
    copies of variables */
#pragma omp parallel private(nthreads, tid)
 /* Obtain thread number */
 tid = omp get thread num();
 printf("Hello from thread = %d\n", tid);
```

```
/* Only master thread does this */
 if (tid == 0)
  nthreads = omp_get_num_threads();
  printf("Number of threads = %d\n", nthreads);
 } /* All threads join master thread and disband */
 n_1 = 1.0 / (double)INTERVALS;
/* Parallel loop with reduction for calculating PI */
#pragma omp parallel for private(i,x)
shared(n_1,INTERVALS) reduction(+:pi)
 for (i = 0; i < INTERVALS; i++)
  x = n 1 * ((double)i - 0.5);
  pi += 4.0 / (1.0 + x * x);
  pi *= n_1;
  printf ("Pi = \%.12lf\n", pi);
```

OpenMP result: 1-D Heat Equation

\$ sbatch -A use300 pi_openmp.sb

```
[mahidhar@login01 openmp_examples]$ more pi.1336540.exp-1-06.out
Hello from thread = 6
Hello from thread = 8
Hello from thread = 11
Threre are 16 threads!
Hello from thread = 4
Hello from thread = 2
Hello from thread = 1
Hello from thread = 0
Hello from thread = 15
Hello from thread = 5
Hello from thread = 13
Hello from thread = 14
Hello from thread = 9
Hello from thread = 3
Hello from thread = 7
Hello from thread = 12
Hello from thread = 10
PI = 3.139384000000, Err = 2.208654e-03, Time = 0.045669
[mahidhar@login01 openmp_examples]$
```



More Work-Share Constructs

- SECTIONS directive enclosed sections are divided among the threads.
- WORKSHARE directive divides execution of block into units of work, each of which is executed once.
- SINGLE directive Enclosed code is executed by only one thread.

Example SECTIONS code

```
!$OMP PARALLEL SHARED(A,B,C,D), PRIVATE(I)
!$OMP SECTIONS
!$OMP SECTION
     DO I = 1, N
         C(I) = A(I) + B(I)
      ENDDO
!$OMP SECTION
     DO I = 1, N
         D(I) = A(I) * B(I)
      ENDDO
!$OMP END SECTIONS NOWAIT
!$OMP END PARALLEL
```

Synchronization Constructs

- MASTER directive Specifies region is executed only by the master thread.
- CRITICAL directive Region of the code that is executed one thread at a time.
- BARRIER directive synchronize all threads
- TASKWAIT directive wait for all child tasks to complete
- ATOMIC directive specific memory location updated atomically (not let all threads write at the same time)



Simple Application using OpenMP: 1-D Heat Equation

- $\partial T/\partial t = \alpha(\partial^2 T/\partial x^2)$; T(0) = 0; T(1) = 0; $(0 \le x \le 1)$ T(x,0) is know as an initial condition.
- Discretizing for numerical solution we get: $T^{(n+1)}_{i} T^{(n)}_{i} = (\alpha \Delta t / \Delta x^{2})(T^{(n)}_{i-1} 2T^{(n)}_{i} + T^{(n)}_{i+1})$ (*n* is the index in time and *i* is the index in space)

Fortran OpenMP Code: 1-D Heat Equation

```
PROGRAM HEATEON
                                                            implicit none
                                                            ************
   integer :: iglobal, itime, nthreads
                                                              pi = 4d0*datan(1d0)
   real*8 :: xalp,delx,delt,pi
                                                              do iglobal = 0, 10
                                                               T(0,iglobal) = dsin(pi*delx*dfloat(iglobal))
   real*8 :: T(0:100,0:10)
                                                              enddo
   integer:: id
                                                           ****** Iterations
   integer:: OMP GET THREAD NUM,
                                                            ****************
    OMP GET NUM THREADS
                                                              do itime = 1.3
                                                               write(*,*)"Running Iteration Number ", itime
!$OMP PARALLEL SHARED(nthreads)
                                                           !$OMP PARALLEL DO PRIVATE(iglobal)
!$OMP MASTER
                                                           SHARED(T,xalp,delx,delt,itime)
   nthreads = omp_get_num_threads()
                                                               do iglobal = 1,9
                                                               T(itime,iglobal)=T(itime-1,iglobal)+
   write (*,*) 'There are', nthreads, 'threads'
                                                              + xalp*delt/delx/delx*
!$OMP END MASTER
                                                              + (T(itime-1,iglobal-1)-2*T(itime-1,iglobal)+T(itime-
!$OMP END PARALLEL
                                                           1,iglobal+1))
   if (nthreads.ne.3) then
                                                               enddo
    write(*,*)"Use exactly 3 threads for this case"
                                                           !SOMP BARRIER
    stop
                                                              enddo
   endif
                                                              do iglobal = 0, 10
                                                               write(*,*)iglobal,T(3, iglobal)
   delx = 0.1d0
                                                              enddo
   delt = 1d-4
                                                              END
   xalp = 2.0d0
```



OpenMP result: 1-D Heat Equation

\$ sbatch heat_openmp.sb

Sample output file:

There a	re 3 threads	
Running	g Iteration Number	•
Running	g Iteration Number	
Running	g Iteration Number	
0	0.00000000000000E+00	0
1	0.307205621017285	
2	0.584339815421976	
3	0.804274757358271	
4	0.945481682332598	
5	0.994138272681972	
6	0.945481682332598	
7	0.804274757358271	
8	0.584339815421977	
9	0.307205621017285	

10 7.797843424221369E-316



Run Time Library Routines

- Setting and querying number of threads
- Querying thread identifier, team size
- Setting and querying dynamic threads feature
- Querying if in parallel region and at what level
- Setting and querying nested parallelism
- Setting, initializing and terminating locks, nested locks.
- Querying wall clock time and resolution.



Environment Variables

- OMP_SCHEDULE e.g set to "dynamic"
- OMP_NUM_THREADS
- OMP_DYNAMIC (TRUE or FALSE)
- OMP_PROC_BIND (TRUE or FALSE)
- OMP_NESTED (TRUE of FALSE)
- OMP_STACKSIZE size of stack for created threads
- OMP_THREAD_LIMIT



General OpenMP Performance Considerations

- Avoid or minimize use of BARRIER, CRITICAL (complete serialization here!), ORDERED regions, and locks. Can use NOWAIT clause to avoid redundant barriers.
- Parallelize at a high level, i.e. maximize the work in the parallel regions to reduce parallelization overhead.
- Use appropriate loop scheduling static has low synchronization overhead but can be unbalanced, dynamic (and guided) have higher synchronization overheads but can improve load balancing.
- Avoid false sharing (more about it in following slide)!



What is False Sharing?

- Most modern processors have a cache buffer between slow memory and high speed registers of the CPU.
- Accessing a memory location causes a "cache line" to be copied into the cache.
- In an OpenMP code two processors may be accessing two different elements in the same cache line. On writes this will lead to "cache line" being marked invalid (because cache coherency is being maintained).
- This will lead to an increase in memory traffic even though the write is to different elements (hence the term false sharing).
- This can have a drastic performance impact if such updates are occurring frequently in a loop.



False Sharing Example

```
Code snippet:
double global=0.0, local[NUM_THREADS];
#pragma omp parallel num_threads(NUM_THREADS)
int tid = omp_get_thread_num();
local[tid] = 0.0;
#pragma omp for
for (i = 0; i < N; i++)
local[tid] += x[i];
#pragma omp atomic
global += local[me];
```

False Sharing - Solutions

Three options

- Compiler directives to align individual variables on cache line boundaries
- Pad arrays/data structures to make sure array elements begin on cache line boundary.
- Use thread local copies of data (assuming the copy overhead is small compared to overall run time).

Homework!

- Download matrix multiply example from LLNL site:
 - https://computing.llnl.gov/tutorials/openMP/samples/Fortran/omp_mm.f
 (wget
 https://computing.llnl.gov/tutorials/openMP/samples/Fortran/omp_mm.f
 on
 Expanse)
- Compile (ifort -fopenmp omp_mm.f) and run the example.
 See if you can vary the environmental variables, scheduling to get better performance!
- This is very quick intro. Lot of ongoing developments.
 Detailed specifications at:
 - https://www.openmp.org/specifications/
- Excellent tutorials from LLNL:
 - https://computing.llnl.gov/tutorials/openMP/

